

# **SLP ASSIGNMENT 1**

**NAME : VARUN D**

**ROLL NO : 717822I262**

**DEPT : ARTIFICIAL INTELLIGENCE  
AND DATA SCIENCE**

**SUBJECT: SPEECH AND LANGUAGE  
PROCESSING**

**COURSE CODE: 21ID31**

# TITLE: Computational Semantics and Semantic Parsing:

## QUESTION:

Build a semantic parser for translating natural language questions into executable queries over a knowledge base or database. Train the parser using a dataset of question-query pairs and evaluate its performance in terms of accuracy and execution speed.

### Procedure for Building a Semantic Parser (NL → SQL)

#### 1. Data Collection:

- Prepare a dataset of natural language questions paired with their corresponding SQL queries (e.g., "What is the population of Canada?" -> SELECT population FROM countries WHERE name = 'Canada';).

#### 2. Model Initialization:

- Use a pre-trained transformer model (like BERT) and tokenizer from Hugging Face. Fine-tune the model using the question-query dataset.

#### 3. Training:

- Create a custom dataset class to tokenize the question-query pairs. Split the dataset into training and validation sets. Use the Trainer API to train the model on the training data.

#### 4. SQL Generation:

- After training, deploy the model with a function that takes a natural language question as input and outputs a generated SQL query based on the model's predictions.

#### 5. Deployment:

- Use Gradio to create a simple web interface where users can input questions and view the corresponding SQL query generated by the model.

## CODE:

### Import necessary libraries:

```
# Importing necessary libraries
from transformers import BertTokenizer, BertForSequenceClassification, Trainer, TrainingArguments
import torch
from torch.utils.data import Dataset
```

### Dataset of questions and corresponding SQL queries

```
# New Sample dataset of questions and corresponding SQL queries
data = [
    {"question": "What is the population of Canada?", "query": "SELECT population FROM countries WHERE name = 'Canada'"},
    {"question": "Which country has the most languages?", "query": "SELECT name FROM countries ORDER BY languages DESC LIMIT 1"},
    {"question": "What is the capital of Italy?", "query": "SELECT capital FROM countries WHERE name = 'Italy'"},
    {"question": "List all countries in Asia.", "query": "SELECT name FROM countries WHERE continent = 'Asia'"}
]
```

```
Downloading nvidia_cublas_cu12-12.4.5.8-py3-none-manylinux2014_x86_64.whl (363.4 MB)
_____ 363.4/363.4 MB 4.1 MB/s eta 0:00:00
Downloading nvidia_cuda_cupti_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (13.8 MB)
_____ 13.8/13.8 MB 25.3 MB/s eta 0:00:00
Downloading nvidia_cuda_nvrtc_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (24.6 MB)
_____ 24.6/24.6 MB 26.7 MB/s eta 0:00:00
Downloading nvidia_cuda_runtime_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (883 kB)
_____ 883.7/883.7 kB 15.7 MB/s eta 0:00:00
Downloading nvidia_cudnn_cu12-9.1.0.70-py3-none-manylinux2014_x86_64.whl (664.8 MB)
_____ 664.8/664.8 MB 2.6 MB/s eta 0:00:00
Downloading nvidia_cufft_cu12-11.2.1.3-py3-none-manylinux2014_x86_64.whl (211.5 MB)
_____ 211.5/211.5 MB 6.2 MB/s eta 0:00:00
Downloading nvidia_curand_cu12-10.3.5.147-py3-none-manylinux2014_x86_64.whl (56.3 MB)
_____ 56.3/56.3 MB 11.7 MB/s eta 0:00:00
Downloading nvidia_cusolver_cu12-11.6.1.9-py3-none-manylinux2014_x86_64.whl (127.9 MB)
_____ 127.9/127.9 MB 7.0 MB/s eta 0:00:00
Downloading nvidia_cusparsparse_cu12-12.3.1.170-py3-none-manylinux2014_x86_64.whl (207.5 MB)
_____ 207.5/207.5 MB 5.4 MB/s eta 0:00:00
Downloading nvidia_nvjitlink_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl (21.1 MB)
_____ 21.1/21.1 MB 78.8 MB/s eta 0:00:00
Downloading datasets-3.2.0-py3-none-any.whl (480 kB)
_____ 480.6/480.6 kB 36.1 MB/s eta 0:00:00
Downloading dill-0.3.8-py3-none-any.whl (116 kB)
_____ 116.3/116.3 kB 10.5 MB/s eta 0:00:00
Downloading fsspec-2024.9.0-py3-none-any.whl (179 kB)
```

## Tokenizer and Model Initialization

### # Dataset class

```
class QuestionQueryDataset(Dataset):
    def __init__(self, data, tokenizer):
        self.data = data
        self.tokenizer = tokenizer

    def __len__(self):
        return len(self.data)

    def __getitem__(self, idx):
        item = self.data[idx]
        question = item['question']
        query = item['query']

        # Tokenizing the question and the query together
        encoding = self.tokenizer(question, query, truncation=True, padding='max_length', max_length=128, return_tensors="pt")
        return {
            'input_ids': encoding['input_ids'].squeeze(0),
            'attention_mask': encoding['attention_mask'].squeeze(0),
            'labels': torch.tensor(0, dtype=torch.float) # Adjust labels if necessary
        }
```

### # Create Train Dataset with new data

```
train_dataset = QuestionQueryDataset(data, tokenizer)
```

### # New validation dataset

```
val_data = [
    {"question": "What is the population of Australia?", "query": "SELECT population FROM countries WHERE name = 'Australia'"},
    {"question": "Which country has the highest GDP?", "query": "SELECT name FROM countries ORDER BY gdp DESC LIMIT 1"}
]

val_dataset = QuestionQueryDataset(val_data, tokenizer)
```

## TRAINING ARGUMENTS

```
training_args = TrainingArguments(
    output_dir="./results",
    evaluation_strategy="epoch",
    learning_rate=2e-5,
    per_device_train_batch_size=4,
    per_device_eval_batch_size=4,
    num_train_epochs=3,
    report_to="none",
)
```

## INITIALIZING TRAINER WITH THE VALIDATION DATASET

# Initialize Trainer with Validation Dataset

```
trainer = Trainer(  
    model=model,  
    args=training_args,  
    train_dataset=train_dataset,  
    eval_dataset=val_dataset,  
)
```

# Start training

```
trainer.train()
```

```
tokenizer_config.json: 100% ██████████ 48.0/48.0 [00:00<00:00, 2.06kB/s]  
vocab.txt: 100% ██████████ 232k/232k [00:00<00:00, 660kB/s]  
tokenizer.json: 100% ██████████ 466k/466k [00:00<00:00, 2.62MB/s]  
config.json: 100% ██████████ 570/570 [00:00<00:00, 35.6kB/s]  
model.safetensors: 100% ██████████ 440M/440M [00:01<00:00, 240MB/s]  
Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized: ['classifier.bias', 'classifier.weight'].  
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.  
/usr/local/lib/python3.11/dist-packages/transformers/training_args.py:1575: FutureWarning: `evaluation_strategy` is deprecated and will be removed in version 4.46 of 🤗 Transformers.  
warnings.warn(  
██████████ [3/3 00:22, Epoch 3/3]
```

Epoch	Training Loss	Validation Loss
1	No log	0.000906
2	No log	0.005842
3	No log	0.006461

```
TrainOutput(global_step=3, training_loss=0.06107855339845022, metrics={'train_runtime': 30.6448, 'train_samples_per_second': 0.392, 'train_steps_per_second': 0.098, 'total_flos': 789326078976.0, 'train_loss': 0.06107855339845022, 'epoch': 3.0})
```

## DEPLOYMENT

```
import gradio as gr
import torch
from transformers import BertTokenizer, BertForSequenceClassification

# Load tokenizer and model (Ensure you have a trained model)
model_name = "bert-base-uncased"
tokenizer = BertTokenizer.from_pretrained(model_name)
model = BertForSequenceClassification.from_pretrained(model_name, num_labels=4) # Adjust num_labels to match your dataset size

# Function to predict SQL queries
def predict_sql(question):
    encoding = tokenizer(question, truncation=True, padding="max_length", max_length=128, return_tensors="pt")

    with torch.no_grad():
        output = model(**encoding)

    # Convert model output (logits) into a meaningful SQL query
    logits = output.logits
    predicted_class = torch.argmax(logits, dim=1).item() # Get the predicted class

    # Example of mapping predicted classes to queries (Customize this!)
    sql_queries = {
        0: "SELECT population FROM countries WHERE name = 'Canada';",
        1: "SELECT name FROM countries ORDER BY languages DESC LIMIT 1;",
        2: "SELECT capital FROM countries WHERE name = 'Italy';",
        3: "SELECT name FROM countries WHERE continent = 'Asia';"
    }

    return sql_queries.get(predicted_class, "Unable to generate SQL query.")

# Create Gradio interface
interface = gr.Interface(
    fn=predict_sql,
    inputs=gr.Textbox(lines=2, placeholder="Type your question here..."),
    outputs=gr.Textbox(label="Generated SQL Command"),
    title="Natural Language to SQL Generator",
    description="Ask a question, and the model will generate the corresponding SQL query."
)

# Launch the interface
interface.launch()
```

\*

GRADIO:

PUBLIC URL: <https://c1a3f9846ea92edaea.gradio.live/>

OUTPUT:

### Natural Language to SQL Generator

Ask a question, and the model will generate the corresponding SQL query.

question

Tell me the population of Canada

Clear

Generated SQL Command

SELECT name FROM countries ORDER BY languages DESC LIMIT 1;

Submit

Flag